

Ecorrection

Place student sticker here

Note:

- During the attendance check a sticker containing a unique code will be put on this exam.
- This code contains a unique number that associates this exam with your registration number.
- This number is printed both next to the code and to the signature field in the attendance check list.

Advanced Programming

Exam: IN1503 / Endterm

Date: Monday 22nd February, 2021

Examiner: Prof. Dr. Hans-Joachim Bungartz

Time: 11:30 – 12:30

Working instructions

- This exam consists of **12 pages** with a total of **3 problems**.
Please make sure now that you received a complete copy of the exam.
- The total amount of achievable credits in this exam is 36 credits.
- Detaching pages from the exam is prohibited.
- Allowed resources:
 - This is an open-book exam. The exam is designed having in mind that you can look at the **course material** whenever you want, but don't forget to keep an eye on the time!
- Often, subproblems are independently solvable, so make sure to try everything.
- **Answers are only accepted if the solution approach is documented.** Give a reason for each answer unless explicitly stated otherwise in the respective subproblem.
- Do not write with red or green colors nor use pencils.
- Do not use comments or notes to write your answers. This will not be visible after submission.
- Do not forget to **save** the annotated PDF file. Verify that the annotations are visible in the submission overview.
- Communication with other people during the examination is strictly prohibited.
- If you run into technical issues, we will be available in the usual lecture BBB room, where you can send us a short private message and we will contact you: <https://bbb.in.tum.de/ger-f3u-4w6>. We cannot answer any topic-related questions. In case of doubt, write your assumptions and continue.

Left room from _____ to _____ / Early submission at _____

Problem 1 Working with `std::vector` (11 credits)

a) What does CMake do? Select only one of the following.

- ☐ It makes a compatibility layer between C and C++.
- ☐ It builds a C++ project.
- ☒ It generates the instructions for a build system.

b) Write a C++ function `selectPrint()` that fulfils the following requirements.

The function shall:

- return nothing,
- have a parameter `vec` which represents a `std::vector` over a templated type `T`, and
- select those entries in `vec` that are larger than zero and print them to the command line.

Do not forget to indicate all necessary include statements.

```
#include <vector>
#include <iostream>

template <typename T>
void selectPrint(std::vector<T> &vec) {
    for (const auto& entry: vec) {
        if (entry > 0) { std::cout << entry << std::endl; }
    }
}
```

alternatively with `for_each()`:

```
#include <algorithm>

template <typename T>
void selectPrint_forEach(std::vector<T> &vec) {
    std::for_each(vec.begin(), vec.end(),
        [](T& x){if (x > 0) { std::cout << x << std::endl; }});
}
```

Points:

- 0.5 for all 2 (or 3) include statements
- 1.0 for template correct (template keyword, brackets, typename keyword, `T`, `std::vector<T>`), 0.5 for almost correct.
- 0.5 for return type correct
- 0.5 for mentioning a for loop (or `for_each`)
- 0.5 for correct loop conditions (iterators in `for_each`)
- 0.5 for correct if (or lambda)
- 0.5 for correct command line output

Passing `vec` by value is also fine.

c) Indicate **two** issues that an incompatible type `T` can cause in the function `selectPrint()` (from b)). Explain.

- A type that does not provide an `operator>()` or is not implicitly convertible to an `int` will make the comparison fail. 1 point for any of the two.
- A type for which there is not implementation of `operator<<()` will make the `std::cout <<` fail. 1 point

d) What option do you have in C++ to specify additional restrictions on the type of T in selectPrint() (from b))? Write one sentence explaining.

0
1

concept allows to specify constraints on types for templates. E.g. here, we should restrict to datatypes supporting the ">" operator.

Points:

- name concept (0.5)
- a valid explanation (0.5)

e) The following code shall find the overall number of entries of an integer value target in the vector of integers vec. Indicate **two** runtime/logical errors: point to the corresponding line numbers, give an argument or description what is wrong there, and how each can be fixed.

0
1
2
3

```
1  int & findNumberOfEntries(std::vector<int> &vec, int target) {  
2      int numberOfFoundEntries = 0;  
3      auto i = vec.begin();  
4      while (i != vec.end()) {  
5          i = std::find(vec.begin(), vec.end(), target);  
6          if (i != vec.end()) {  
7              numberOfFoundEntries++;  
8              i++;  
9          }  
10     }  
11     return numberOfFoundEntries;  
12 }
```

1. error: lines 1 and 2: local stack variable reference returned (segfault when accessed outside the function) ⇒ use **int** instead of **int &** as return type
2. error: line 6: loop will not terminate due to start of find always at beginning of vector ⇒ use **i** instead of **vec.begin()** as starting iterator

Points:

- correct errors/lines: 2×0.5
- correct explanation of problems: 2×0.5
- correct proposed solutions: 2×0.5

Problem 2 Object-oriented programming (16 credits)

a) Which of the following do we need to achieve runtime polymorphism?

Check all that apply. A wrong "check" removes a point, with the minimum number of points being zero.

- ☐ A sliced object of a derived class **This is actually what we are trying to avoid by using pointers.**
- ☐ Friend classes to allow access to private members **This is not related to runtime polymorphism**
- ☒ A derived object managed through a pointer to a base class **1 point**
- ☒ Virtual functions **1 point**
- ☐ A virtual constructor **There is actually no such thing as a virtual constructor.**

b) Consider the following code, which implements a Database class to keep track of people that need an appointment for vaccinations:

```
1 struct Person{
2     std::string name;
3     std::size_t age; // age in years
4 };
5
6 class Database{
7 private:
8     Person* _people;
9     std::size_t _num_people;
10    const std::string _author; // Institute that maintains the database, e.g. "RKI".
11
12 public:
13     // (nothing here at the moment)
14 };
```

Write a constructor for Database, which should create a complete and valid state of the object from a given author and num_people and allocates the _people array so that it can store num_people elements. You do not need to give values to the elements of _people.

```
Database(std::string author, std::size_t num_people)
: _author(author), _num_people(num_people), _people(new Person[_num_people])
{}

```

- 0.5 points for using the member initializer list for **const** `std::string _author`.
- 0.5 points for initializing all other members correctly.

c) Write the destructor of Database (from b)) in a way that applies the concept of Resource Acquisition Is Initialization.

```
~Database(){
    delete [] _people;
}

```

- 0.5 points for `~Database()` and **delete**.
- 0.5 points for `[]`.

d) The institute maintains several databases and often wants to construct a new database as a copy of an old one. Implement the copy constructor of Database (from b)).

0
1
2
3

```
// Example implementation:
Database(const Database& from) : _author(from._author){
    _num_people = from._num_people;
    _people = new Person[_num_people];
    // copy people from "from"
    for(auto i = 0; i < _num_people; i++){
        _people[i] = from._people[i];
    }

    // alternatively, use std::copy
}
```

- 1 point for (const) reference `const Database& from`.
- 0.5 points for using the member initializer list `_author(from._author)`.
- 0.5 points for initializing every other member correctly.
- 1 point for copying existing objects from `from._people`, e.g. using a loop or `std::copy`. For almost correct syntax, 0.5 points instead.

e) While developing Database, you quickly realize that you should better rely on existing containers to store your data inside the Database class.

Modify the declaration of `_people` (from b)) so that it is a `std::vector` instead of a pointer.

0
1

```
// Instead of: Person* _people;
std::vector<Person> _people;
```

- 1 point for correct member declaration `std::vector<Person> _people;`. No points for missing the template arguments `<Person>`.

0 ☐

1 ☐

2 ☐

f) Modify the constructor (from b)) and destructor (from c)) accordingly to fit the changes in e).
(Skip the copy constructor – in the following, assume that you also modified the copy constructor here.)

```
// Instead of the previous constructor:
Database(std::string author, std::size_t num_people)
: _author(author), _people(std::vector<Person>(num_people))
{
    _people = std::vector<Person>(num_people);
    //_people.reserve(num_people);
}

// Instead of the previous destructor:
~Database() = default;
// ~Database(){};
// no declaration for destructor necessary
```

- 1 point for correct initialization for member variables in constructor, or 0.5 point for _people of wrong size e.g. not changing _people in constructor
- 1 point for empty ~Database() or for =default or explicitly mentioning that there is no need to declare a destructor.

0 ☐

1 ☐

g) Implement a getter function get_people in Database (from e)) that returns the vector of _people and can be called from outside the class.

```
public:
    std::vector<Person> get_people() {
        return _people;
    }
```

1 point for completely correct getter, 0.5 for missing public, 0 points otherwise.
Returning by reference is also fine. Returning by const reference is accepted, even though it creates issues later.

0 ☐

1 ☐

h) In main(), construct a new_database from the old_database using only functions/methods defined in the class Database (from f)). In case you skipped the previous parts: we are now using vectors and you can assume the copy constructor defined.

```
1 int main(){
2     Database old_database(...); // Assume given
3
4 }
```

```
Database old_database(...); // Assume given
Database new_database(old_database);
// Alternative: Database new_database = old_database;
// Wrong: Database new_database(...);
//         new_database = old_database;
```

1 point for calling the copy constructor to create new_database from the given old_database.
0 points for using the copy-assignment operator.

i) Consider the following algorithm, which rearranges a container based on a condition:

```
template< class ForwardIt, class UnaryPredicate >
ForwardIt partition( ForwardIt first, ForwardIt last, UnaryPredicate p )
```

where the predicate is a function that returns `true` or `false` for each element pointed to by the iterators (in the range `[first, last)`).

Get the vector of people from `new_database` (from `h`) using `get_people`. Using the `std::partition` algorithm, partition the vector into two parts: people with age greater than or equal to 65 and people with age less than 65. The order of the two parts does not matter in this case.

```
auto vec = new_database.get_people();
auto mid = std::partition(vec.begin(), vec.end(),
                        [](Person p) { return p.age >= 65;});
// alternative without assigning to mid ok.
// alternative using std::begin: std::partition(std::begin(vec), ...)
// alternative to lambda: using the function bool pred(Person p){...}
```

- 1 point for using the `.begin()` and `.end()` in `std::partition`:
`std::partition(vec.begin(), vec.end(), ...);`
- 1 point for the general syntax `[] (Person p) {}` OR `bool pred(Person p) {}`.
- 1 point for correct body of function: `{return p.age >= 65;}` or `{return p.age < 65;}`.
0.5 points for slight mistakes, e.g. `{return p.age > 65;}` or `{return p.age <= 65;}` (no ;).

j) Given a `const` database, for example:

```
const Database new_database(...);
```

what are the requirements so that the following code compiles? Provide specific code changes, if changes are needed.


```
auto vec = new_database.get_people();
```

The `get_people()` method needs to be `const`-correct:

```
std::vector<Person> get_people() const;
```

- 1 point for correct function signature `std::vector<Person> get_people() const;`
- 0 points for vaguely mentioning "const" or putting it in the wrong side.

Problem 3 Performance analysis, optimization, and vectorization (9 credits)



a) Consider the following code kernel:

```
1  for(auto i = 0; i < points.size(); i++){
2      auto elem = points[i];
3      result[i] = 8 + elem + 2 * elem * elem;
4  }
```

where points and result are of type `std::vector<double>` and of size N.

You are ordering a new computer and you have the choice between two processors with the following differences:

Model A: vector units that can perform 8 double-precision FLOP/cycle.

Model B: vector units that can perform 16 double-precision FLOP/cycle.

Both processors have the same frequency (2GHz) and scalar performance (1 FLOP/cycle), while “Model B” is significantly more expensive. In both cases, the memory bandwidth of the system will be 12 GB/s (same for read- and write-operations).

Which processor would you buy, with the only application being the above code kernel? Explain your decision thoroughly using the roofline model analysis.

Memory interactions of this code: 1 load + 1 store for each iterations, $\rightarrow 2 \cdot N$ memory interactions.

Total floating point operations: $4 \cdot N$ (2 additions, 2 multiplications).

Computational intensity: $(4 \cdot N \text{ FLOP}) / (2 \cdot 8 \cdot N \text{ Byte}) = 1/4 \text{ FLOP/Byte}$. Remember: one double needs 8 Byte and we have two memory interactions.

Theoretically maximum performance that the code can achieve: $1/4 \text{ FLOP/Byte} \cdot 12 \text{ GByte/s} = 3 \text{ GFLOP/s}$.

Peak scalar performance for both systems: $1 \text{ FLOP/cycle} \cdot 2 \text{ Gcycle/s} = 2 \text{ GFLOP/s}$ (the code would be compute-bound).

The loop can be vectorized. Vector peak performance for “Model A”: $8 \text{ FLOP/cycle} \cdot 2 \text{ GHz} = 16 \text{ GFLOP/s}$. Vector peak performance for “Model B”: $16 \text{ FLOP/cycle} \cdot 2 \text{ GHz} = 32 \text{ GFLOP/s}$.

Since the theoretical maximum performance that the given code can achieve with this memory bandwidth is below the vector performance of both processors (memory-bound), we decide to get the cheaper “Model A”.

- 1 point for correct intensity (including units).
Only 0.5 points for not dividing by 8 Bytes or for forgetting units.
- 1 point for correct peak performance of the code.
- 1 point for correct computation of the peak vector performance of the two models (the scalar performance computation is optional).
- 1 point for showing that it is memory-bound for the vector case and deciding for “Model A”.

0
1

b) Is the following loop vectorizable? Explain.

```
1 // double result[N];
2 // double arr[N];
3
4 for(auto i = 0; i < N-4; i++){
5     result[i] += arr[i];
6     result[i+1] += arr[i];
7     result[i+2] += arr[i];
8     result[i+3] += arr[i];
9 }
```

No: there are loop-carried dependencies (e.g. in `result[i+1]`). If the loop step was `i+4`, this would not be an issue. 1 point for correct explanation, 0 points otherwise.

0
1

c) Is the following loop vectorizable? Explain.

```
1 // double result[N];
2 // double arr[N];
3
4 for(auto i = 1; i < N; i++){
5     result[i] = 2 * arr[i-1];
6 }
```

Yes: Straight-line of code with a known number of iterations. 1 point for correct explanation, 0 points otherwise.

0
1
2

d) Consider the following code:

```
1 //res initialized to zero
2 // double res[];
3 // double mat[][];
4 for(auto col= 0; col < N; col++){
5     for(auto row = 0; row < N; row++){
6         res[row] += mat[row][col] / 2;
7     }
8 }
```

Explain at least two modifications to improve cache efficiency and/or computation costs (assuming no compiler optimizations).

- Modification for cache efficiency: swap the order of loops:

```
for(auto row = 0; row < N; row++){
    for(auto col= 0; col < N; col++){
        res[row] += mat[row][col] / 2;
    }
}
```

Cache lines follow the matrix rows in C++.

- Modification for reducing the computation costs:

```
res[row] += mat[row][col] * 0.5;
```

Multiplication is cheaper than division.

1 point for each correct argumentation. Different optimizations accepted if valid and explained.

0 ☐ e) You are developing a simulation program, which needs different digits of pi. The number of needed digits is known at compile time and the compute_pi function is (only) called once.

1 ☐

```
1     double compute_pi(int num_digits){
2         // performs expensive computation
3     }
4     int main(){
5         double pi = compute_pi(8);
6         // use throughout program
7     }
```

Is there any way for you to optimize the runtime behavior of this code? Show any code changes needed.

Yes, we can define both compute_pi and pi as **constexpr**:

```
constexpr double compute_pi(int num_digits){
    // performs computation
}
int main(){
    constexpr double pi = compute_pi(8);
    // use throughout program
}
```

- 0.5 points for **constexpr** function
- 0.5 points for **constexpr** double pi

Additional space for solutions—clearly mark the (sub)problem your answers are related to and strike out invalid solutions.



Sample Solution

Correction Notes